PROGRAM 1

OBJECTIVE: Write a program to implement Logistic Regression.

CODE:

import libraries and dataset

import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

df=pd.read_csv("insurance_data.csv")

#data Preprocessing

df.isnull().sum()

df.shape

df.describe()

#split the data

x=df.iloc[:,:1]

y=df.iloc[:,1:]

from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.4,random_state=23)

#train the model

from sklearn.linear_model import LogisticRegression

model=LogisticRegression()

model.fit(x_train,y_train)

#model prediction

y_pred=model.predict(x_test)

#finding Evaluation matrics

from sklearn.metrics import confusion_matrix

import matplotlib.pyplot as plt

```
import seaborn as sns

cf=confusion_matrix(y_test,y_pred)

plt.figure()

sns.heatmap(cf,annot=True)

plt.xlabel('Prediction')

plt.ylabel('Target')

plt.title('Confusion matrix')

from sklearn.metrics import accuracy_score

accuracy= accuracy_score(y_test,y_pred)

print("accuracy:{:.2f}%".format(accuracy*100))

from sklearn.metrics import precision_score,recall_score,f1_score

print("Precision score:{:.2f}%".format(precision_score(y_test,y_pred)*100))

print("Recall Score::{:.2f}%".format(recall_score(y_test,y_pred)*100))

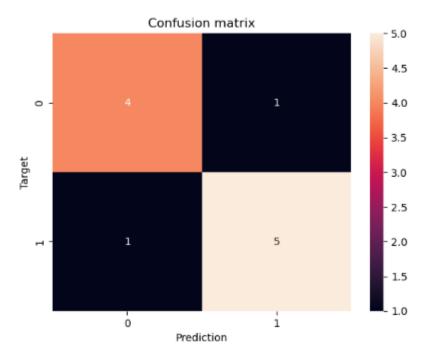
print("f1 score:{:.2f}%".format(f1_score(y_test,y_pred)*100))
```

```
In [2]: df=pd.read_csv("insurance_data.csv")
Out[2]:
              age bought_insurance
           0 22
           2
               47
               52
                                0
           3
           4
               48
           5 56
           6 55
               60
In [6]: x=df.iloc[:,:1]
           y=df.iloc[:,1:]
 In [7]: x
Out[7]:
                age
                 22
                 25
             2
                 47
                 52
             4
                 46
In [8]: y
Out[8]:
               bought_insurance
            0
                               0
            1
            2
            3
                               0
            5
                               1
 In [12]: from sklearn.model_selection import train_test_split
            x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.4,random_state=23)
In [15]: from sklearn.linear_model import LogisticRegression
         model=LogisticRegression()
In [16]: model.fit(x_train,y_train)
         C:\Users\hp\AppData\Local\anaconda3\Lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector
         y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().

y = column_or_1d(y, warn=True)
Out[16]: - LogisticRegression
         LogisticRegression()
In [17]: y_pred=model.predict(x_test)
```

```
In [18]: from sklearn.metrics import confusion_matrix
    import matplotlib.pyplot as plt
    import seaborn as sns
    cf=confusion_matrix(y_test,y_pred)
    plt.figure()
    sns.heatmap(cf,annot=True)
    plt.xlabel('Prediction')
    plt.ylabel('Target')
    plt.title('Confusion matrix')
```

Out[18]: Text(0.5, 1.0, 'Confusion matrix')



In [20]: from sklearn.metrics import precision_score,recall_score,f1_score
 print("Precision score:{:.2f}%".format(precision_score(y_test,y_pred)*100))
 print("Recall Score:{:.2f}%".format(recall_score(y_test,y_pred)*100))
 print("f1 score:{:.2f}%".format(f1_score(y_test,y_pred)*100))

Precision score:83.33% Recall Score::83.33% f1 score:83.33%

PROGRAM 2

OBJECTIVE: Write a program to implement Linear Regression with one variables.

CODE:

```
# import libraries
```

import numpy as np

import pandas as pd

from sklearn.model selection import train test split

from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_squared_error

import matplotlib.pyplot as plt

Load the dataset

```
data = {'Years_of_Experience': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10], 'Salary': [40000, 45000, 50000, 55000, 60000, 65000, 70000, 75000, 80000, 85000]}
```

df = pd.DataFrame(data)

Split the data

```
X = df['Years\_of\_Experience'].values.reshape(-1, 1)
```

y = df['Salary'].values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,random_state=42)

Train the model

model = LinearRegression()

model.fit(X_train, y_train)

Make predictions

y_pred = model.predict(X_test)

Evaluate the model

```
mse = mean_squared_error(y_test, y_pred)
```

print(f"Mean Squared Error: {mse}")

Plotting

```
plt.scatter(X, y, color='blue')
plt.plot(X, model.predict(X), color='red')
plt.title('Years of Experience vs. Salary')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.show()
```

```
In [16]: #inport libraries for signle variable linear regression
         import numpy as np
         import pandas as pd
         from sklearn.model_selection import train_test_split
         from sklearn.linear_model import LinearRegression
         from sklearn.metrics import mean_squared_error
         import matplotlib.pyplot as plt
In [17]: # Load the dataset
         data = {'Years_of_Experience': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10], 'Salary':
                 [40000, 45000, 50000, 55000, 60000, 65000, 70000, 75000, 80000, 85000]}
         df = pd.DataFrame(data)
In [18]: # Split the data
         X = df['Years_of_Experience'].values.reshape(-1, 1)
         y = df['Salary'].values
         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                              random_state=42)
In [19]: # Train the model
         model = LinearRegression()
         model.fit(X_train, y_train)
Out[19]: + LinearRegression
         LinearRegression()
In [20]: # Make predictions
         y_pred = model.predict(X_test)
         # Evaluate the model
         mse = mean_squared_error(y_test, y_pred)
         print(f"Mean Squared Error: {mse}")
        Mean Squared Error: 0.0
In [21]: # Plotting
         plt.scatter(X, y, color='blue')
         plt.plot(X, model.predict(X), color='red')
        plt.title('Years of Experience vs. Salary')
         plt.xlabel('Years of Experience')
        plt.ylabel('Salary')
         plt.show()
                                   Years of Experience vs. Salary
           80000
           70000
        Salary
00009
           50000
           40000
                                                                    8
                                                                                 10
                                          Years of Experience
```

PROGRAM 3

OBJECTIVE: Write a program to implement Linear Regression with two variables.

```
CODE:
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
df=pd.read_csv("boston-housing-dataset.csv")
df.isnull().sum()
df.shape
df.describe()
x = df.iloc[:,0:13] #diving data into x and y
X
y=df.iloc[:,13]
y
from sklearn.model_selection import train_test_split
                                                        #dividing data into train and test based on
randomstate method
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.10,random_state=40)
from sklearn.preprocessing import StandardScaler
scaler =StandardScaler()
                            # to call only standardize function and not to train(just call)
scaler.fit_transform(x_train) # fit is to give the standard scaler data to transform(standardize) and fit
intox train data
X train =scaler.fit_transform(x_train) # only train
X_test = scaler.transform(x_test) # only change
from sklearn.linear_model import LinearRegression
#cross validation
from sklearn.model_selection import cross_val_score
# estimator
```

regression = LinearRegression()

regression.fit(X_train,y_train)

cross_val_score(regression,X_train,y_train,scoring='neg_mean_squared_error',cv=10) #these
valueswill be different for all systems as values are taken random so score will also vary

mse=cross_val_score(regression,X_train,y_train,scoring='neg_mean_squared_error',cv=10)

#storing values into mse

np.mean(mse)

#prediction

reg_pred =regression.predict(X_test)

reg_pred

import seaborn as sns

sns.displot(reg_pred-y_test)

from sklearn.metrics import r2_score

score =r2_score(reg_pred,y_test)

score

```
In [4]: df=pd.read_csv("boston-housing-dataset.csv")
Out[4]:
               Unnamed: 0
                             CRIM ZN INDUS CHAS NOX
                                                               RM AGE
                                                                            DIS RAD TAX PTRATIO
                                                                                                          B LSTAT MEDV
                        0 0.00632 18.0
                                                                    65.2 4.0900
                                                                                    1 296.0
                                                                                                 15.3 396.90
                                                                                                                      24.0
                                                    0 0.538 6.575
                                                                                                               4.98
                                           7.07
                                                    0 0.469 6.421 78.9 4.9671
                                                                                   2 242.0
            1
                        1 0.02731
                                    0.0
                                                                                                 17.8 396.90
                                                                                                               9.14
                                                                                                                      21.6
            2
                        2 0.02729 0.0
                                          7.07
                                                    0 0.469 7.185 61.1 4.9671
                                                                                   2 242.0
                                                                                                 17.8 392.83
                                                                                                               4.03
                                                                                                                      34.7
            3
                         3 0.03237
                                   0.0
                                          2.18
                                                    0 0.458 6.998 45.8 6.0622
                                                                                   3 222.0
                                                                                                 18.7 394.63
                                                                                                               2.94
                                                                                                                      33.4
                         4 0.06905 0.0
                                          2.18
                                                    0 0.458 7.147 54.2 6.0822
                                                                                   3 222 0
                                                                                                 18.7 396.90
                                                                                                               5.33
                                                                                                                      38.2
 In [6]: x= df.iloc[:,0:13] #diving data into x and y
 In [7]: y=df.iloc[:,13]
 In [8]: from sklearn.model_selection import train_test_split
                                                                   #dividing data into train and test based on random state method
         x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.10,random_state=40)
 In [9]: x_train # but here the decimal values are discrete(1,4,7 etc)
 out[9]:
              Unnamed: 0
                           CRIM ZN INDUS CHAS NOX RM AGE
                                                                    DIS RAD TAX PTRATIO
          12
                     12 0.09378 12.5
                                       7.87
                                               0 0.524 5.889 39.0 5.4509
                                                                           5 311.0
                                                                                       15.2 390.50
          310
                     310 2.63548 0.0
                                      9.90
                                               0 0.544 4.973 37.8 2.5194
                                                                           4 304.0
                                                                                       18.4 350.45
          405
                     405 67.92080 0.0 18.10
                                            0 0.693 5.683 100.0 1.4254
                                                                          24 666.0
                                                                                      20.2 384.97
          418
                     418 73.53410 0.0 18.10
                                               0 0.879 5.957 100.0 1.8028
                                                                          24 666.0
                                                                                       20.2 16.45
                                                                                   14.7 388.45
                    163 1.51902 0.0 19.58
                                            1 0.605 8.375 93.9 2.1620
          163
                                                                           5 403.0
          440
                     440 22.05110 0.0 18.10
                                            0 0.740 5.818 92.4 1.8662
                                                                          24 666.0
                                                                                      20.2 391.45
          165
                     185 2.92400 0.0 19.58
                                               0 0.805 8.101 93.0 2.2834
                                                                           5 403.0
                                                                                       14.7 240.18
                     7 0.14455 12.5 7.87
                                            0 0.524 8.172 98.1 5.9505
                                                                                       15.2 396.90
          219
                     219 0.11425 0.0 13.89
                                               1 0.550 6.373 92.4 3.3633
                                                                           5 278.0
                                                                                       18.4 393.74
                     326 0.30347 0.0 7.38
                                            0 0.493 8.312 28.9 5.4159
                                                                                      19.6 396.90
          326
                                                                         5 287.0
         455 rows × 13 columns
In [10]: from sklearn.preprocessing import StandardScaler
                                   # to call only standardize function and not to train(just call)
         scaler =StandardScaler()
 In [11]: scaler.fit_transform(x_train) # fit is to give the standard scaler data to transform(standardize) and fit into x_train_data
 Out[11]: array([[-1.66314065, -0.39493589, 0.05881085, ..., -0.56712089,
                  -1.54500291, 0.37941884],
                 [ 0.37212357, -0.09175343, -0.48418707, ..., -0.60880171,
                  -0.03697567, -0.05089125],
                 [ 1.02094942, 7.69569747, -0.48418707, ..., 1.54669214,
                   0.81128966, 0.32000274],
                 [-1.69728938, -0.38887987, 0.05881085, ..., -0.56712089,
                  -1.54500291, 0.4481825 ],
                 [-0.24938329, -0.39249416, -0.48418707, ..., -0.77552499,
                 -0.9794927 , 0.41423044],
[ 0.48139951, -0.36992336, -0.48418707, ..., -0.71002656,
                  0.52853455, 0.4481825 ]])
 In [12]: X_train =scaler.fit_transform(x_train) # only train
          X_test = scaler.transform(x_test) # only change
```

```
In [15]: from sklearn.linear_model import LinearRegression
          #cross validation
          from sklearn.model_selection import cross_val_score
In [16]: # estimator
         regression = LinearRegression()
         regression.fit(x_train,y_train)
Out[16]: TinearRegression
          LinearRegression()
In [17]: cross_val_score(regression,X_train,y_train,scoring='neg_mean_squared_error',cv=10) #these values will be differ
Out[17]: array([-28.94179834, -23.0608436 , -15.54632857, -18.0865388 , -8.95692519, -17.01503615, -11.73314584, -22.88137448,
                 -11.60561642, -17.09032995])
In [18]: mse=cross_val_score(regression,X_train,y_train,scoring='neg_mean_squared_error',cv=10) #storing values into mse
In [19]: import numpy as no
         np.mean(mse)
Out[19]: -17.49179373207918
In [20]: #prediction
         reg_pred =regression.predict(X_test)
In [21]: reg_pred
Out[21]: array([16.07879814, 5.00069276, 16.95394955, 18.0443779 , 15.4976683 ,
                 14.14550016, 8.63256373, 16.38754087, 16.70900917, 29.4034721,
                 11.21496112, 15.92697665, 22.65808116, 24.76002238, 10.35786763,
                  9.4243188 , 12.07721013, 16.19729959, 13.148523 , 17.11632489,
                  1.98682757, 20.38472578, 3.60728592, 7.97657611, 2.02481831,
                  3.87492122, 8.60224097, 15.70259083, 15.22821993, 14.39847259,
                  9.48291381, 11.47115589, 9.12462972, 9.23418362, 26.28911803,
                  6.11357086, 6.39458915, 20.14564815, 3.97699318, 8.59140618,
                 11.13966363, 21.5752782 , 15.01162204, 11.24311788, 12.30292074,
                 21.5319964 , 17.4820148 , 5.35171674, 22.02567034, 7.38005371, 15.98160939])
 In [22]: import seaborn as sns
         sns.displot(reg_pred-y_test)
Out[22]: <seaborn.axisgrid.FacetGrid at 0x1cae0a01f50>
             14
             12
             10
              8
              6
              2
                -15
                      -10
                              -5
                                      0
In [23]: from sklearn.metrics import r2_score
In [24]: score =r2_score(reg_pred,y_test)
In [25]: score
Out[25]: 0.2287798024913632
```

PROGRAM 4

OBJECTIVE: Write a program to implement the naïve Bayesian classifier for a sample training data set stored as a .CSV file. Compute the accuracy of the classifier, considering few test data sets.

CODE:

```
import pandas as pd
from sklearn import tree
from sklearn.preprocessing import LabelEncoder
from sklearn.naive_bayes import GaussianNB
data = pd.read_csv('tennis.csv')
print("The first 5 values of data is:\n",data.head())
#splitting data
X = data.iloc[:,:-1]
print("\nThe First 5 values of train data is\n",X.head())
y = data.iloc[:,-1]
print("\nThe first 5 values of Train outpu tis\n",y.head())
# Convert then in numbers
le_outlook = LabelEncoder()
X.outlook = le outlook.fit transform(X.outlook)
le_temp = LabelEncoder()
X.temp = le_temp.fit_transform(X.temp)
le humidity = LabelEncoder()
X.humidity =le_humidity.fit_transform(X.humidity)
le_windy = LabelEncoder()
X.windy =le_windy.fit_transform(X.windy)
print("\nNow the Train data is :\n",X.head())
le_play = LabelEncoder()
```

```
y = le_play.fit_transform(y)
print("\nNow the Train output is\n",y)
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test =train_test_split(X,y, test_size=0.20)
classifier = GaussianNB()
classifier.fit(X_train,y_train)
from sklearn.metrics import accuracy_score
print("Accuracy is:",accuracy_score(classifier.predict(X_test),y_test))
```

```
In [1]: import pandas as pd
          from sklearn import tree
          from sklearn.preprocessing import LabelEncoder
          from sklearn.naive_bayes import GaussianNB
In [3]: data = pd.read_csv('tennis.csv')
print("THe first 5 values of data is:\n",data.head())
         THe first 5 values of data is:
               outlook temp humidity windy play
sunny hot high False no
                        hot
                                        True
               sunny
                                 high
         2 overcast hot
                                 high False yes
               rainy mild
                               high False yes
               rainy cool normal False yes
In [4]: #splitting data
         X = data.iloc[:,:-1]
         print("\nThe First 5 values of train datais\n",X.head())
         y = data.iloc[:,-1]
         print("\nThe first 5 values of Train outputis\n",y.head())
         The First 5 values of train datais
              outlook temp humidity windy
               sunny
                        hot
                                 high False
               sunny
                                 high
         2 overcast hot
                                 high False
                              high False
             rainy mild
         4
               rainy cool normal False
         The first 5 values of Train outputis
          0
               no
         1
                no
               ves
In [8]: # Convert then in numbers
          le outlook = LabelEncoder()
           X.outlook =le_outlook.fit_transform(X.outlook)
          le_temp = LabelEncoder()
          X.temp =le_temp.fit_transform(X.temp)
          le_humidity = LabelEncoder()
X.humidity =le_humidity.fit_transform(X.humidity)
le_windy = LabelEncoder()
          X.windy =le_windy.fit_transform(X.windy)
print("\nNow the Train data is :\n",X.head())
le_play = LabelEncoder()
          y = le_play.fit_transform(y)
          print("\nNow the Train output is\n",y)
          Now the Train data is :
               outlook temp humidity windy
                                       0
          2
                     0
                            1
                                       ø
                                               0
          3
                    1
                           2
                                       ø
                                               0
                           0
                    1
          Now the Train output is
           [00111010111110]
In [9]: from sklearn.model_selection import train_test_split
          X_train, X_test, y_train, y_test =train_test_split(X,y, test_size=0.20)
classifier = GaussianNB()
classifier.fit(X_train,y_train)
Out[9]: - GaussianNB
           GaussianNB()
In [10]: from sklearn.metrics import accuracy_score
          print("Accuracy is:",accuracy_score(classifier.predict(X_test),y_test))
          Accuracy is: 0.666666666666666
```

PROGRAM 5

OBJECTIVE: Implement a K-Nearest Neighbors (KNN) classifier from scratch in Python. Use a sample dataset, such as the Iris dataset, and split it into a training and testing set. Train the KNN classifier on the training set and evaluate its performance on the testing set. Experiment with different values of k and report the accuracy of the classifier. CODE:

```
#importing libraries
import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
#Model creation from scratch
def most_common(lst):
  return max(set(lst), key=lst.count)
def euclidean(point, data):
  # Euclidean distance between points a & data
  return np.sqrt(np.sum((point - data)**2, axis=1))
class KNeighborsClassifier:
  def___init_(self, k=5, dist_metric=euclidean):
    self.k = k
     self.dist_metric = dist_metric
  def fit(self, X_train, y_train):
     self.X_train = X_train
     self.y_train = y_train
  def predict(self, X_test):
    neighbors = []
```

for x in X_test:

```
distances = self.dist_metric(x, self.X_train)
       y_sorted = [y for _, y in sorted(zip(distances, self.y_train))]
       neighbors.append(y_sorted[:self.k])
     return list(map(most_common, neighbors))
  def evaluate(self, X_test, y_test):
     y_pred = self.predict(X_test)
     accuracy = sum(y_pred == y_test) / len(y_test)
     return accuracy
# unpack the iris dataset.
iris = datasets.load_iris()
X = iris['data']
y = iris['target']
# Split data into train & test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20)
# Preprocessing of dataset
ss = StandardScaler().fit(X_train)
X_{train}, X_{test} = ss.transform(X_{train}), ss.transform(X_{test})
# Test knn model across varying ks
accuracies = []
ks = range(1, 30)
for k in ks:
  knn = KNeighborsClassifier(k=k)
  knn.fit(X_train, y_train)
  accuracy = knn.evaluate(X_test, y_test)
  accuracies.append(accuracy)
#outputs
```

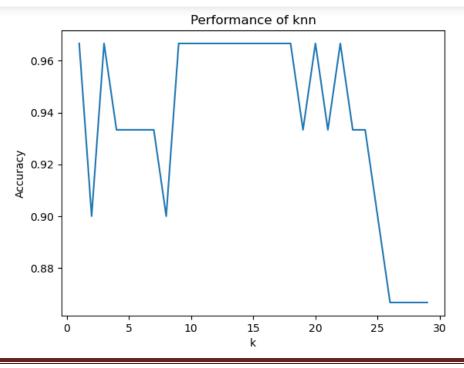
```
from sklearn.metrics import classification_report,confusion_matrix,accuracy_score
accuracy=accuracy_score(y_test,knn.predict(X_test))
print("Accuracy: {:.2f}%".format(accuracy * 100))
print("Confusion Matrix:\n",confusion_matrix(y_test,knn.predict(X_test)))
print("Classification Report:\n",classification_report(y_test,knn.predict(X_test)))
# Visualize accuracy vs. k
fig, ax = plt.subplots()
ax.plot(ks, accuracies)
ax.set(xlabel="k",
    ylabel="Accuracy",
    title="Performance of knn")
plt.show()
```

```
#importing necessary libraries
import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```
#Model creation from scratch
def most_common(lst):
    return max(set(lst), key=lst.count)
def euclidean(point, data):
    # Euclidean distance between points a & data
    return np.sqrt(np.sum((point - data)**2, axis=1))
class KNeighborsClassifier:
    def __init__(self, k=5, dist_metric=euclidean):
        self.k = k
        self.dist_metric = dist_metric
    def fit(self, X train, y train):
        self.X_train = X_train
        self.y_train = y_train
    def predict(self, X_test):
        neighbors = []
        for x in X test:
            distances = self.dist_metric(x, self.X_train)
             y_sorted = [y for _, y in sorted(zip(distances, self.y_train))]
             neighbors.append(y_sorted[:self.k])
        return list(map(most_common, neighbors))
    def evaluate(self, X_test, y_test):
        y_pred = self.predict(X_test)
        accuracy = sum(y_pred == y_test) / len(y_test)
        return accuracy
: # Unpack the iris dataset, from UCI Machine Learning Repository
 iris = datasets.load_iris()
 X = iris['data']
 y = iris['target']
```

```
[5. , 3.4, 1.5, 0.2],
     [4.4, 2.9, 1.4, 0.2],
     [4.9, 3.1, 1.5, 0.1],
     [5.4, 3.7, 1.5, 0.2],
     [4.8, 3.4, 1.6, 0.2],
     [4.8, 3., 1.4, 0.1],
     [4.3, 3. , 1.1, 0.1],
: y
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
    : # Split data into train & test sets
 X_train, X_test, y_train, y_test =train_test_split(X, y, test_size=0.20)
: # Preprocessing of dataset
 ss = StandardScaler().fit(X_train)
 X_train, X_test = ss.transform(X_train),ss.transform(X_test)
```

```
# Test knn model across varying ks
  accuracies = []
 ks = range(1, 30)
  for k in ks:
      knn = KNeighborsClassifier(k=k)
      knn.fit(X_train, y_train)
      accuracy = knn.evaluate(X test, y test)
      accuracies.append(accuracy)
: #outputs
  from sklearn.metrics import classification_report,confusion_matrix,accuracy_score
  accuracy=accuracy_score(y_test,knn.predict(X_test))
  print("Accuracy: {:.2f}%".format(accuracy * 100))
  print("Confusion Matrix:\n",confusion_matrix(y_test,knn.predict(X_test)))
  print("Classification Report:\n",classification_report(y_test,knn.predict(X_test)))
  Accuracy: 86.67%
  Confusion Matrix:
  [[10 0 0]
  [ 0 8 2]
[ 0 2 8]]
  Classification Report:
               precision
                            recall f1-score
                                             support
            0
                   1.00
                             1.00
                                      1.00
                                                 10
                   0.80
                             0.80
                                      0.80
            1
                                                 10
            2
                   0.80
                             0.80
                                      0.80
                                                 10
                                      0.87
                                                 30
     accuracy
    macro avg
                   0.87
                             0.87
                                      0.87
                                                 30
  weighted avg
                   0.87
                             0.87
                                      0.87
  # Visualize accuracy vs. k
  fig, ax = plt.subplots()
  ax.plot(ks, accuracies)
  ax.set(xlabel="k",
          ylabel="Accuracy",
          title="Performance of knn")
  plt.show()
```



PROGRAM 6

OBJECTIVE: Write a python program to implement support Vector Machine(SVM) classifier using a library like scikit-learn. choose a suitable dataset for binary classification(e.g., the Breast Cancer dataset) and split it into training and testing sets. Train the SVM classifier on the training data and evaluate its performance on the testing data, reporting metrics such as accuracy, precision, recall, and F1-score. Experiment with different kernel functions (e.g., linear, radial basis function) and compare their performance.

CODE:

#importing libraries

```
import numpy as np
```

import matplotlib.pyplot as plt

from sklearn import datasets

from sklearn.model_selection import train_test_split

from sklearn.svm import SVC

from sklearn.metrics import accuracy_score,precision_score, recall_score, f1_score,confusion_matrix

Loading the dataset

```
cancer = datasets.load_breast_cancer()
```

X = cancer.data

y = cancer.target

Splitting the dataset into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,random_state=42)
```

#Model creation

```
def evaluate_classifier(clf, X_test, y_test):
```

```
y_pred = clf.predict(X_test)
```

accuracy = accuracy_score(y_test,y_pred)

precision = precision_score(y_test,y_pred)

recall = recall_score(y_test, y_pred)

f1 = f1_score(y_test, y_pred)

confusion_mat =confusion_matrix(y_test, y_pred)

return accuracy, precision, recall, f1,confusion_mat

```
# Try SVM with different kernel functions
```

```
kernel_functions = ['linear','rbf'] for kernel in kernel_functions:
```

Create and train the SVM classifier

```
svm_classifier = SVC(kernel=kernel)
svm_classifier.fit(X_train, y_train)
```

Evaluate the classifier

```
accuracy, precision, recall, f1,confusion_mat =evaluate_classifier(svm_classifier, X_test, y_test)
```

#results

```
print(f'\nResults for SVM with {kernel} kernel:')
```

#Output:

```
print(f'Accuracy: {accuracy:.4f}')
print(f'Precision: {precision:.4f}')
print(f'Recall: {recall:.4f}')
print(f'F1-score: {f1:.4f}')
print(f'Confusion Matrix:\n{confusion_mat}')
```

```
import numpy as np
 import matplotlib.pyplot as plt
 from sklearn import datasets
 from sklearn.model selection import train test split
 from sklearn.svm import SVC
 from sklearn.metrics import accuracy_score,precision_score, recall_score, f1_score,confusion_matrix
 # Loading the dataset
 cancer = datasets.load_breast_cancer()
 X = cancer.data
 y = cancer.target
 array([[1.799e+01, 1.038e+01, 1.228e+02, ..., 2.654e-01, 4.601e-01,
           1.189e-01],
          [2.057e+01, 1.777e+01, 1.329e+02, ..., 1.860e-01, 2.750e-01,
           8.902e-02],
 0, 0, 0,
1, 0, 0,
                            0, 0, 1, 0, 1,
                                            0, 0, 1, 1, 1, 0, 0,
         1, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1,
         1, 1, 1, 1, 1, 1, 0, 0, 0,
                                         1, 0, 0, 1, 1, 1, 0, 0,
         0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1,
         1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0,
 X_train, X_test, y_train, y_test =train_test_split(X, y, test_size=0.2,random_state=42)
 #Model creation
 def evaluate_classifier(clf, X_test, y_test):
    y_pred = clf.predict(X_test)
     precision = precision_score(y_test,y_pred)
     recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
confusion_mat = confusion_matrix(y_test, y_pred)
 return accuracy, precision, recall, f1,confusion_mat
# Try SVM with different kernel functions
kernel_functions = ['linear','rbf']
 for kernel in kernel_functions:
 # Create and train the SVM classifier
svm classifier = SVC(kernel=kernel)
      svm_classifier.fit(X_train, y_train)
 accuracy, precision, recall, f1,confusion mat =evaluate classifier(svm classifier, X test, y test)
 print(f'\nResults for SVM with {kernel} kernel:')
 #Output:
 print(f'Accuracy: {accuracy:.4f}')
print(f'Precision: {precision:.4f}')
 print(f'Recall: {recall:.4f}')
 print(f'F1-score: {f1:.4f}')
 print(f'Confusion Matrix:\n{confusion_mat}')
 Results for SVM with rbf kernel:
 Accuracy: 0.9474
 Precision: 0.9221
 Recall: 1.0000
 F1-score: 0.9595
 Confusion Matrix:
 [[37 6]
[ 0 71]]
: #results
print(f'\nResults for SVM with {kernel} kernel:')
  print(f'Accuracy: {accuracy:.4f}')
print(f'Precision: {precision:.4f}')
print(f'Recall: {recall:.4f}')
print(f'F1-score: {f1:.4f}')
  print(f'Confusion Matrix:\n{confusion_mat}')
  Results for SVM with linear kernel:
  Accuracy: 0.9561
  Precision: 0.9459
  Recall: 0.9859
  F1-score: 0.9655
  Confusion Matrix:
  [[39 4]
[170]]
```

PROGRAM 7

OBJECTIVE: Given a dataset containing features and labels, implement a Random Forest classification model using Python and a library like scikit-learn. Split the dataset into training and testing sets, train the model, and evaluate its performance using metrics like accuracy, precision, and recall.

CODE:

#importing libraries

import numpy as np

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score, precision_score, recall_score, confusion_matrix

from sklearn.datasets import load_iris

#loading dataset

```
iris.data = load_iris()
```

X = iris.data

y = iris.target

X

y

Splitting the dataset into training and testing sets

```
X_train, X_test, y_train, y_test =train_test_split(X, y, test_size=0.30,random_state=42)
```

Creating a Random Forest Classifier

rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)

Training the model

rf_classifier.fit(X_train, y_train)

Making predictions on test set

y_pred = rf_classifier.predict(X_test)

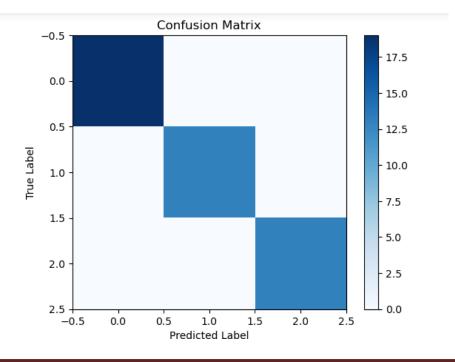
Evaluating the classifier

```
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
conf_mat = confusion_matrix(y_test, y_pred)
# Printing performance metrics
print(f'Accuracy: {accuracy:.4f}')
print(f'Precision: {precision:.4f}')
print(f'Recall: {recall:.4f}')
print('Confusion Matrix:')
print(conf_mat)
# Plotting the confusion matrix
plt.imshow(conf_mat, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.colorbar()
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```

```
#importing necessary libraries
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, confusion_matrix
from sklearn.datasets import load_iris
# Loading the dataset
iris = load iris()
X = iris.data
y = iris.target
array([[5.1, 3.5, 1.4, 0.2],
     [4.9, 3., 1.4, 0.2],
     [4.7, 3.2, 1.3, 0.2],
     [4.6, 3.1, 1.5, 0.2],
у
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
     # Splitting the dataset into training and testing sets
X_train, X_test, y_train, y_test =train_test_split(X, y, test_size=0.30,random_state=42)
# Creating a Random Forest Classifier
rf classifier = RandomForestClassifier(n_estimators=100, random_state=42)
# Training the model
rf_classifier.fit(X_train, y_train)
       RandomForestClassifier
RandomForestClassifier(random_state=42)
```

```
# Making predictions on the test set
y_pred = rf_classifier.predict(X_test)
# Evaluating the classifier
accuracy = accuracy score(y test, y pred)
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
conf_mat = confusion_matrix(y_test, y_pred)
# Printing performance metrics
print(f'Accuracy: {accuracy:.4f}')
print(f'Precision: {precision:.4f}')
print(f'Recall: {recall:.4f}')
print('Confusion Matrix:')
print(conf_mat)
Accuracy: 1.0000
Precision: 1.0000
Recall: 1.0000
Confusion Matrix:
[[19 0 0]
[ 0 13 0]
 [0 0 13]]
```

```
# Plotting the confusion matrix
plt.imshow(conf_mat, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.colorbar()
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()
```



PROGRAM 8

OBJECTIVE: Design a Hebb Net to implement logical AND function. CODE:

```
#ImportingLibrary.
```

import numpy as np

Define the input patterns and target for the AND function

```
input_patterns = np.array([[1, 1], [1, -1], [-1, 1], [-1, -1]])
target_output = np.array([1, -1, -1, -1])
weights = np.zeros(input_patterns.shape[1])
bias = 0
weights
```

Train the Hebb Net

```
for i in range(input_patterns.shape[0]):
    weights += input_patterns[i] * target_output[i]
    bias += target_output[i]
    print(f"weights{i}: ", weights, f"\nbias{i}: ", bias)
```

Define function to make predictions with trained Hebb Net

```
def predict(input_pattern):
    return 1 if np.dot(input_pattern, weights) + bias >= 0 else 0
```

Test the trained Hebb Net

```
for input_pattern in input_patterns:
```

```
print(f"Input: {input_pattern}, Output: {predict(input_pattern)}")
```

```
In [21]: #importing numpy
         import numpy as np
In [22]: # Define the input patterns and target output for the AND function
         input_patterns = np.array([[1,1], [1,-1], [-1, 1], [-1, -1]])
         target_output = np.array([1, -1, -1, -1])
In [23]:
         weights = np.zeros(input_patterns.shape[1])
         bias = 0
         weights
Out[23]: array([0., 0.])
In [25]: # Train the Hebb Net
         for i in range(input_patterns.shape[0]):
            weights += input_patterns[i] * target_output[i]
            bias += target_output[i]
            print(f"weights{i}: ", weights, f"\nbias{i}: ", bias)
       weights0: [3. 3.]
       bias0: -1
       weights1: [2. 4.]
       bias1: -2
       weights2: [3. 3.]
       bias2: -3
       weights3: [4. 4.]
       bias3: -4
          ......
 In [26]: # Define a function to make predictions with the trained Hebb Net
           def predict(input_pattern):
               return 1 if np.dot(input_pattern, weights) + bias >= 0 else 0
 In [27]: # Test the trained Hebb Net
           for input pattern in input patterns:
                print(f"Input: {input pattern}, Output: {predict(input pattern)}")
          Input: [1 1], Output: 1
          Input: [ 1 -1], Output: 0
          Input: [-1 1], Output: 0
          Input: [-1 -1], Output: 0
```

PROGRAM 9

OBJECTIVE: Apply k-Means algorithm to cluster a set of data stored in a .CSV file. Use the same data set for clustering using the Agglomerative Clustering algorithm. Compare the results of these two algorithms and comment on the quality of clustering. You can add Python ML library classes in the program.

CODE:

#import libraries

import numpy as np

import pandas as pd

from sklearn.cluster import KMeans, AgglomerativeClustering

import matplotlib.pyplot as plt

from sklearn.metrics import silhouette_score

from sklearn.datasets import load_iris

Loading dataset

iris = load iris()

X = iris.data

we are using the first two features for clustering

X = X[:, :2]

No need to scale the features in this

#Apply k-Means clustering

kmeans = KMeans(n_clusters=3,random_state=42)

kmeans_labels = kmeans.fit_predict(X)

Apply hierarchical clustering forcomparison

hierarchical = AgglomerativeClustering(n_clusters=3)

hierarchical_labels = hierarchical.fit_predict(X)

Visualize the results

plt.figure(figsize=(12, 5))

```
plt.subplot(1, 2, 1)
plt.scatter(X[:, 0], X[:, 1], c=kmeans_labels,cmap='viridis', edgecolors='k', s=50)
plt.title('k-Means Clustering')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.subplot(1, 2, 2)
plt.scatter(X[:, 0], X[:, 1],
c=hierarchical_labels, cmap='viridis',edgecolors='k', s=50)
plt.title('Hierarchical Clustering')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.tight_layout()
plt.show()
# Compare clustering quality using silhouette score
kmeans_silhouette = silhouette_score(X,kmeans_labels)
#Output
hierarchical_silhouette = silhouette_score(X,hierarchical_labels)
print(f'Silhouette Score - k-Means:{kmeans_silhouette:.4f}')
print(f'Silhouette Score - Hierarchical:{hierarchical_silhouette:.4f}')
```

OUTPUT:

#importing libraries

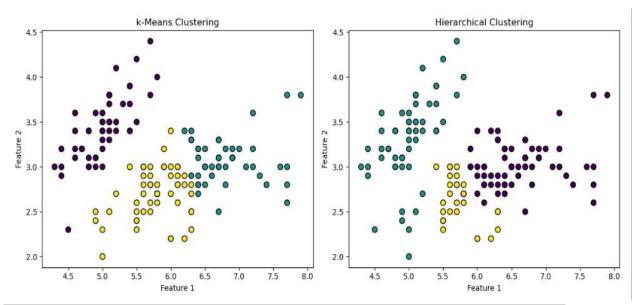
```
import numpy as np
import pandas as pd
from sklearn.cluster import KMeans, AgglomerativeClustering
import matplotlib.pyplot as plt
from sklearn.metrics import silhouette score
from sklearn.datasets import load_iris
# Loading the dataset
iris = load iris()
X = iris.data
array([[5.1, 3.5, 1.4, 0.2],
       [4.9, 3., 1.4, 0.2],
       [4.7, 3.2, 1.3, 0.2],
       [4.6, 3.1, 1.5, 0.2],
       [5., 3.6, 1.4, 0.2],
       [5.4, 3.9, 1.7, 0.4],
       [4.6, 3.4, 1.4, 0.3],
       [5., 3.4, 1.5, 0.2],
       [4.4, 2.9, 1.4, 0.2],
       [4.9, 3.1, 1.5, 0.1],
       [5.4, 3.7, 1.5, 0.2],
# Assume we are using the first two features for clustering
X = X[:, :2]
# No need to scale the features in this case
# Apply k-Means clustering
kmeans = KMeans(n clusters=3,random state=42)
kmeans_labels = kmeans.fit_predict(X)
# Apply hierarchical clustering forcomparison
hierarchical =AgglomerativeClustering(n_clusters=3)
hierarchical_labels =hierarchical.fit_predict(X)
# Visualize the results
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.scatter(X[:,\ 0],\ X[:,\ 1],\ c=kmeans\_labels,cmap='viridis',\ edgecolors='k',\ s=50)
plt.title('k-Means Clustering')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.subplot(1, 2, 2)
plt.scatter(X[:, 0], X[:, 1],
```

c=hierarchical_labels, cmap='viridis',edgecolors='k', s=50)

plt.title('Hierarchical Clustering')

plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.tight_layout()

plt.show()



```
# Compare clustering quality using silhouette score
kmeans_silhouette = silhouette_score(X,kmeans_labels)
#Output
hierarchical_silhouette = silhouette_score(X,hierarchical_labels)
print(f'Silhouette Score - k-Means:{kmeans_silhouette:.4f}')
print(f'Silhouette Score - Hierarchical:{hierarchical_silhouette:.4f}')
Silhouette Score - k-Means:0.4451
Silhouette Score - Hierarchical:0.3653
```



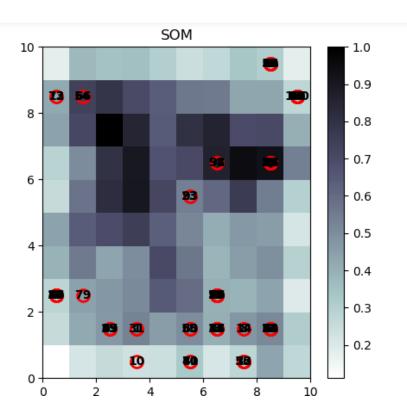
OBJECTIVE: Write a Python program to implement a Self-Organizing Map (SOM) and train it on a given dataset, such as a collection of 2D points. Allow the user to specify parameters like the map size, learning rate, and number of training iterations. Visualize the map before and after training to observe how it adapts to the data.

CODE:

```
#numpy based SOM implimentation
!pip install minisom
#importing libraries
import numpy as np
from minisom import MiniSom
import matplotlib.pyplot as plt
#for visualization
def visualize_som(som, data, title):
  plt.figure(figsize=(5, 5))
  plt.pcolor(som.distance_map().T, cmap='bone_r') # plot the distance map as background
  plt.colorbar()
 #for visualization
def visualize_som(som, data, title):
  plt.figure(figsize=(5, 5))
  plt.pcolor(som.distance_map().T, cmap='bone_r') # plot the distance map as background
  plt.colorbar()
  # plot points on the map
  for i(x, \underline{\ }) in enumerate(data):
     w = som.winner(x)
    plt.plot(w[0] + 0.5, w[1] + 0.5, 'o', markerfacecolor='None', markersize=10, markeredgecolor='r',
    markeredgewidth=2)
    plt.text(w[0] + 0.5, w[1] + 0.5, str(i + 1), color='k', fontweight='bold', ha='center', va='center')
```

```
plt.title("SOM")
  plt.show()
# Generate synthetic 2D data
np.random.seed(42)
data = np.random.rand(100, 2) # replace this with your own dataset
# User-defined parameters
map\_size = (10, 10) # SOM map size
learning_rate = 0.5 # initial learning rate
num_iterations = 1000 # number of training iterations
# Visualize the SOM before training
visualize_som(som, data, title="SOM Before Training")
# Create and train the SOM
som = MiniSom(*map_size, 2, sigma=1.0, learning_rate=learning_rate)
som.random_weights_init(data)
print("Training SOM...")
som.train_random(data, num_iterations)
print("Training complete.")
# Visualize the SOM after training
visualize_som(som, data, title="SOM After Training")
```

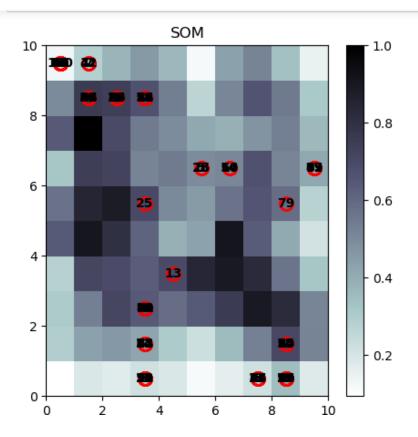
```
In [5]: #numpy based SOM implimentation
           !pip install minison
           Collecting minison
             Downloading MiniSom-2.3.1.tar.gz (10 kB)
             Preparing metadata (setup.py): started
             Preparing metadata (setup.py): finished with status 'done'
           Building wheels for collected packages: minisom
Building wheel for minisom (setup.py): started
             Building wheel for minisom (setup.py): finished with status 'done'
             Created wheel for minisom: filename=MiniSom-2.3.1-py3-none-any.whl size=10601 sha256=36bf28d62436e67a3444b7eec0218017a0b772e0
           bbd305fe764135a52ea1517d
             Stored\ in\ directory:\ c:\users\hp\appdata\local\pip\cache\wheels\28\e3\3d\7877393fa9813d5ab7b3ffb914ded8ca3c48dec231fa392528
           Successfully built minison
           Installing collected packages: minisom
           Successfully installed minisom-2.3.1
: #importing necessary libraries
  import numpy as np
  from minisom import MiniSom
  import matplotlib.pyplot as plt
: #for visualization
  def visualize_som(som, data, title):
      plt.figure(figsize=(5, 5))
      plt.pcolor(som.distance_map().T, cmap='bone_r') # plot the distance map as background
      plt.colorbar()
      # plot data points on the map
      for i, (x, _) in enumerate(data):
          w = som.winner(x)
          plt.plot(w[0] + 0.5, w[1] + 0.5, 'o', markerfacecolor='None', markersize=10, markeredgecolor='r', markeredgewidth=2)
          plt.text(w[0] + 0.5, w[1] + 0.5, str(i + 1), color='k', fontweight='bold', ha='center', va='center')
      plt.title("SOM")
      plt.show()
  # Generate synthetic 2D data
  np.random.seed(42)
  data = np.random.rand(100, 2) # replace this with your own dataset
  # User-defined parameters
  map_size = (10, 10) # SOM map size
  learning rate = 0.5 # initial learning rate
  num iterations = 1000 # number of training iterations
  # Visualize the SOM before training
  visualize_som(som, data, title="SOM Before Training")
```



```
# Create and train the SOM
som = MiniSom(*map_size, 2, sigma=1.0, learning_rate=learning_rate)
som.random_weights_init(data)
print("Training SOM...")
som.train_random(data, num_iterations)
print("Training complete.")
```

Training SOM...
Training complete.

```
# Visualize the SOM after training
visualize_som(som, data, title="SOM After Training")
```



PROGRAM 11

OBJECTIVE: (A) Take a binary classification dataset and implement both the K-Nearest Neighbors (KNN) and Support Vector Machine (SVM) classifiers using Python. Compare the performance of these two algorithms on metrics such as accuracy, precision, recall, and F1-score. Visualize the decision boundaries for both algorithms.

CODE:

#installing mlxtend

pip install mlxtend

#importing necessary libraries

import numpy as np

import matplotlib.pyplot as plt

from sklearn import datasets

from sklearn.model_selection import train_test_split

from sklearn.neighbors import KNeighborsClassifier

from sklearn.svm import SVC

from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix

from mlxtend.plotting import plot_decision_regions

Loading Iris dataset (for binary classification)

iris = datasets.load iris()

X = iris.data[:, :2] # Selecting only the first two features for visualization

y = (iris.target != 0).astype(int) # Convert to binary classification

Splitting the dataset into training and testing sets

X_train, X_test, y_train, y_test =train_test_split(X, y, test_size=0.20,random_state=32)

K-Nearest Neighbors (KNN) Classifier

knn_classifier = KNeighborsClassifier(n_neighbors=3)

knn_classifier.fit(X_train, y_train)

y_pred_knn = knn_classifier.predict(X_test)

Support Vector Machine (SVM) Classifier

```
svm_classifier = SVC(kernel='linear')
svm_classifier.fit(X_train, y_train)
y pred svm = svm classifier.predict(X test)#
Evaluate classifiers
def evaluate_classifier(y_true, y_pred):
  accuracy = accuracy_score(y_true, y_pred)
  precision = precision_score(y_true, y_pred)
  recall = recall_score(y_true, y_pred)
  f1 = f1_score(y_true, y_pred)
  confusion_mat = confusion_matrix(y_true, y_pred)
  return accuracy, precision, recall, f1,confusion_mat
# Evaluate KNN classifier
accuracy_knn, precision_knn, recall_knn,fl_knn, confusion_mat_knn = evaluate_classifier(y_test,
y_pred_knn)
# Evaluate SVM classifier
accuracy sym, precision sym, recall sym, f1 sym, confusion mat sym = evaluate classifier(y test,
y_pred_svm)
# Visualize decision boundaries
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
fig.suptitle('Decision Boundaries of KNN and SVM')
# Decision boundary for KNN
plot_decision_regions(X, y, clf=knn_classifier, legend=2, ax=axes[0])
axes[0].set_title('K-Nearest Neighbors (KNN)')
# Decision boundary for SVM
plot_decision_regions(X, y,
clf=svm_classifier, legend=2, ax=axes[1])
```

```
axes[1].set_title('Support Vector Machine(SVM)')

plt.show()

# Print performance metrics

print('\nPerformance Metrics for K-Nearest Neighbors (KNN):')

print(f'Accuracy: {accuracy_knn:.4f}')

print(f'Precision: {precision_knn:.4f}')

print(f'Recall: {recall_knn:.4f}')

print(f'Confusion Matrix:\n{confusion_mat_knn}')

print(f'Accuracy: {accuracy_svm:.4f}')

print(f'Accuracy: {accuracy_svm:.4f}')

print(f'Precision: {precision_svm:.4f}')

print(f'Recall: {recall_svm:.4f}')

print(f'F1-score: {f1_svm:.4f}')

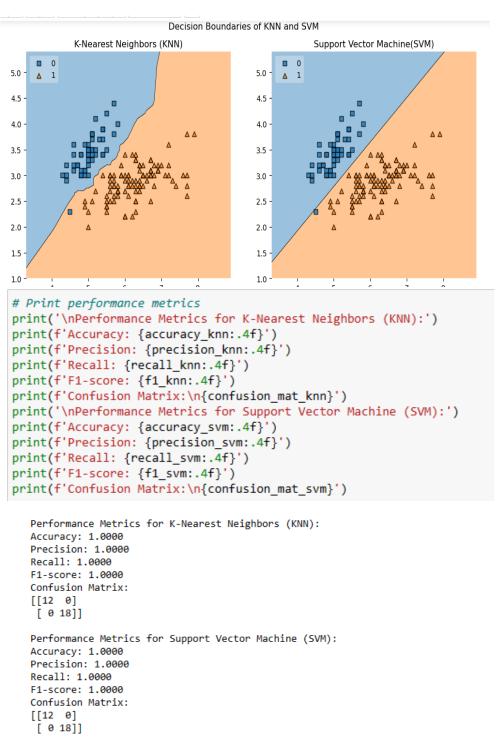
print(f'F1-score: {f1_svm:.4f}')

print(f'Confusion Matrix:\n{confusion_mat_svm}')
```

OUTPUT:

```
#installing mlxtend
pip install mlxtend
Collecting mlxtend
  Downloading mlxtend-0.23.0-py3-none-any.whl (1.4 MB)
                                                  0.0/1.4 MB ? eta -:--:--
                                                  0.2/1.4 MB 5.9 MB/s eta 0:00:01
#importing necessary libraries
import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix
from mlxtend.plotting import plot_decision_regions
# Loading Iris dataset (for binary classification)
iris = datasets.load_iris()
X = iris.data[:, :2] # Selecting only the first two features for visualization
y = (iris.target != 0).astype(int) # Convert to binary classification
# Splitting the dataset into training and testing sets
X_train, X_test, y_train, y_test =train_test_split(X, y, test_size=0.20,random_state=32)
# K-Nearest Neighbors (KNN) Classifier
knn_classifier = KNeighborsClassifier(n_neighbors=3)
knn classifier.fit(X train, y train)
y_pred_knn = knn_classifier.predict(X_test)
# Support Vector Machine (SVM) Classifier
svm_classifier = SVC(kernel='linear')
svm_classifier.fit(X_train, y_train)
y_pred_svm = svm_classifier.predict(X_test)
# Evaluate classifiers
def evaluate_classifier(y_true, y_pred):
    accuracy = accuracy_score(y_true, y_pred)
    precision = precision_score(y_true, y_pred)
    recall = recall_score(y_true, y_pred)
    f1 = f1_score(y_true, y_pred)
    confusion_mat = confusion_matrix(y_true, y_pred)
    return accuracy, precision, recall, f1,confusion_mat
# Evaluate KNN classifier
accuracy\_knn, \ precision\_knn, \ recall\_knn, f1\_knn, \ confusion\_mat\_knn = evaluate\_classifier(y\_test, \ y\_pred\_knn)
# Evaluate SVM classifier
accuracy_svm, precision_svm, recall_svm,f1_svm, confusion_mat_svm = evaluate_classifier(y_test, y_pred_svm)
```

```
# Visualize decision boundaries
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
fig.suptitle('Decision Boundaries of KNN and SVM')
# Decision boundary for KNN
plot_decision_regions(X, y, clf=knn_classifier, legend=2, ax=axes[0])
axes[0].set_title('K-Nearest Neighbors (KNN)')
# Decision boundary for SVM
plot_decision_regions(X, y,
clf=svm_classifier, legend=2, ax=axes[1])
axes[1].set_title('Support Vector Machine(SVM)')
plt.show()
```



B) Given a dataset of customer churn, implement a program that compares the performance of three different supervised learning algorithms (e.g., Logistic Regression, Random Forest, and Support Vector Machine) for binary classification. Split the dataset into training and testing sets, train each algorithm on the training set, and evaluate their performance using metrics like accuracy, precision, recall, and F1-score. Present the results in a clear and informative way, such as through a bar chart or a table.

CODE:

#importing necessary libraries import pandas as pd from sklearn.model selection import train test split from sklearn.linear_model import LogisticRegression from sklearn.ensemble import RandomForestClassifier from sklearn.svm import SVC from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score import matplotlib.pyplot as plt # Load your dataset dataset_path = 'telecom_churn.csv' data = pd.read_csv(dataset_path) data.head() # Assuming your dataset has a 'Churn' column indicating binary labels (1 for churn, 0 for nonchurn) X = data.drop('Churn', axis=1)y = data['Churn'] X y # Split the dataset into training and testing sets X_train, X_test, y_train, y_test =train_test_split(X, y, test_size=0.20,random_state=32) **# Define the classifiers**

classifiers = {'Logistic Regression': LogisticRegression(), 'Random Forest':

RandomForestClassifier(), 'Support Vector Machine': SVC()}

```
# Train and evaluate each classifier
results = {'Classifier': [], 'Accuracy': [], 'Precision': [], 'Recall': [], 'F1-Score': []}
for clf name, clf in classifiers.items():
  # Train the classifier
  clf.fit(X train, y train)#
  Predict on the test set
  y_pred = clf.predict(X_test)
  # Evaluate performance
  accuracy = accuracy_score(y_test, y_pred)
  precision = precision_score(y_test, y_pred)
  recall = recall_score(y_test, y_pred)
  f1 = f1\_score(y\_test, y\_pred)
  # Store results
  results['Classifier'].append(clf_name)
  results['Accuracy'].append(accuracy)
  results['Precision'].append(precision)
  results['Recall'].append(recall)
  results['F1-Score'].append(f1)
# Convert results to a DataFrame for easy visualization
results_df = pd.DataFrame(results)
# Plot the results using a bar chart
plt.figure(figsize=(7, 5))
for metric in ['Accuracy', 'Precision', 'Recall', 'F1-Score']:
  plt.bar(results_df['Classifier'],results_df[metric], label=metric)
  plt.title('Performance Comparison of Classifiers')
  plt.xlabel('Classifier')
```

```
plt.ylabel('Score')
plt.legend()
plt.show()
# Display the results in tabular form
print("Results:")
print(results_df)
```

OUTPUT:

```
#importing necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score,precision_score, recall_score, f1_score
import matplotlib.pyplot as plt

# Load your dataset
dataset_path = 'telecom_churn.csv'
data = pd.read_csv(dataset_path)
data.head()
```

	Churn	AccountWeeks	ContractRenewal	DataPlan	DataUsage	CustServCalls	DayMins	DayCalls	MonthlyCharge	OverageFee	RoamMins
0	0	128	1	1	2.7	1	265.1	110	89.0	9.87	10.0
1	0	107	1	1	3.7	1	161.6	123	82.0	9.78	13.7
2	0	137	1	0	0.0	0	243.4	114	52.0	6.06	12.2
3	0	84	0	0	0.0	2	299.4	71	57.0	3.10	6.6
4	0	75	0	0	0.0	3	166.7	113	41.0	7.42	10.1

```
# Assuming your dataset has a 'Churn' column indicating binary labels (1 for churn, 0 for nonchurn)
X = data.drop('Churn', axis=1)
y = data['Churn']
X
```

	AccountWeeks	ContractRenewal	DataPlan	DataUsage	CustServCalls	DayMins	DayCalls	MonthlyCharge	OverageFee	RoamMins
0	128	1	1	2.70	1	265.1	110	89.0	9.87	10.0
1	107	1	1	3.70	1	161.6	123	82.0	9.78	13.7
2	137	1	0	0.00	0	243.4	114	52.0	6.06	12.2
3	84	0	0	0.00	2	299.4	71	57.0	3.10	6.6
4	75	0	0	0.00	3	166.7	113	41.0	7.42	10.1
3328	192	1	1	2.67	2	156.2	77	71.7	10.78	9.9
3329	68	1	0	0.34	3	231.1	57	56.4	7.67	9.6
3330	28	1	0	0.00	2	180.8	109	56.0	14.44	14.1
3331	184	0	0	0.00	2	213.8	105	50.0	7.98	5.0
3332	74	1	1	3.70	0	234.4	113	100.0	13.30	13.7

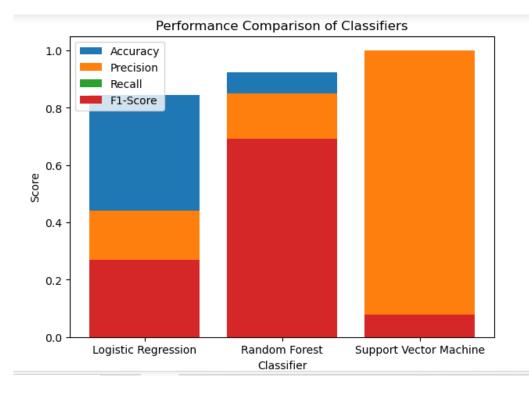
3333 rows x 10 columns

```
: # Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20,random_state=32)
```

```
# Train and evaluate each classifier
results = {'Classifier': [], 'Accuracy': [], 'Precision': [], 'Recall': [], 'F1-Score': []}
for clf name, clf in classifiers.items():
   # Train the classifier
   clf.fit(X_train, y_train)
   # Predict on the test set
   y_pred = clf.predict(X_test)
   # Evaluate performance
   accuracy = accuracy_score(y_test, y_pred)
   precision = precision_score(y_test, y_pred)
   recall = recall_score(y_test, y_pred)
   f1 = f1_score(y_test, y_pred)
   # Store results
   results['Classifier'].append(clf_name)
   results['Accuracy'].append(accuracy)
   results['Precision'].append(precision)
   results['Recall'].append(recall)
   results['F1-Score'].append(f1)
```

```
# Convert results to a DataFrame for easy visualization
results_df = pd.DataFrame(results)
```

```
# Plot the results using a bar chart
plt.figure(figsize=(7, 5))
for metric in ['Accuracy', 'Precision', 'Recall', 'F1-Score']:
    plt.bar(results_df['Classifier'], results_df[metric], label=metric)
    plt.title('Performance Comparison of Classifiers')
    plt.xlabel('Classifier')
    plt.ylabel('Score')
    plt.legend()
plt.show()
```





OBJECTIVE: Write a Python program that loads a dataset and performs an empirical comparison of three different clustering algorithms, such as K-Means, Hierarchical Agglomerative Clustering, and DBSCAN. Evaluate and compare their performance in terms of cluster quality metrics like Silhouette Score or Inertia, and visualize the results.

CODE:

#importing libraries

import pandas as pd

import matplotlib.pyplot as plt

from sklearn import datasets

from sklearn.cluster import KMeans, Agglomerative Clustering, DBSCAN

from sklearn.metrics import silhouette_score

from sklearn.preprocessing import StandardScaler

from sklearn.decomposition import PCA

from sklearn.pipeline import make_pipeline

Load the dataset

iris = datasets.load_iris()

X = iris.data

y = iris.target

X

Standardize the features

scaler = StandardScaler()

 $X_{std} = scaler.fit_{transform}(X)$

Apply PCA for visualization purposes (2D plot)

pca = PCA(n_components=2)

 $X_pca = pca.fit_transform(X_std)$

Define clustering algorithms

kmeans = KMeans(n_clusters=3, random_state=42)

```
hierarchical = AgglomerativeClustering(n_clusters=3)
dbscan = DBSCAN(eps=0.8, min_samples=5)
algorithms = [kmeans, hierarchical, dbscan]
algorithm_names = [ 'K-Means', 'Hierarchical Agglomerative', 'DBSCAN']
# Evaluate and visualize each clustering algorithm
for algorithm, algorithm_name in zip(algorithms, algorithm_names):
# Fit the clustering algorithm to the data
  if algorithm_name != 'DBSCAN':
     algorithm.fit(X_std)
    labels=algorithm.labels_
  else:
    labels = algorithm.fit\_predict(X\_std)
  # Evaluate clustering quality using Silhouette Score
  silhouette_avg = silhouette_score(X_std, labels)
  print(f"{algorithm_name} Silhouette Score: {silhouette_avg:.4f}")
# Visualize the clustering results in a 2D
plotplt.figure(figsize=(8, 5))
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=labels, cmap='viridis', edgecolor='k', s=50)
plt.title(f'{algorithm_name} Clustering Results (Silhouette Score: {silhouette_avg:.4f})')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()
```

OUTPUT:

```
#importing libraries
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.cluster import KMeans,AgglomerativeClustering, DBSCAN
from sklearn.metrics import silhouette score
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.pipeline import make pipeline
# Load the dataset
iris = datasets.load iris()
X = iris.data
y = iris.target
array([[5.1, 3.5, 1.4, 0.2],
       [4.9, 3., 1.4, 0.2],
      [4.7, 3.2, 1.3, 0.2],
      [4.6, 3.1, 1.5, 0.2],
      [5., 3.6, 1.4, 0.2],
      [5.4, 3.9, 1.7, 0.4],
      [4.6, 3.4, 1.4, 0.3],
      [5., 3.4, 1.5, 0.2],
      [4.4, 2.9, 1.4, 0.2],
      [4 9 3 1 1 5 9 1]
# Standardize the features
scaler = StandardScaler()
X std = scaler.fit transform(X)
# Apply PCA for visualization purposes (2D plot)
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_std)
# Define clustering algorithms
kmeans = KMeans(n_clusters=3, random_state=42)
hierarchical = AgglomerativeClustering(n_clusters=3)
dbscan = DBSCAN(eps=0.8, min_samples=5)
algorithms = [kmeans, hierarchical, dbscan]
algorithm_names = [ 'K-Means', 'Hierarchical Agglomerative', 'DBSCAN']
# Evaluate and visualize each clustering algorithm
for algorithm, algorithm_name in zip(algorithms, algorithm_names):
# Fit the clustering algorithm to the data
    if algorithm_name != 'DBSCAN':
        algorithm.fit(X std)
        labels=algorithm.labels
   else:
        labels = algorithm.fit_predict(X_std)
    # Evaluate clustering quality using Silhouette Score
    silhouette_avg = silhouette_score(X_std, labels)
    print(f"{algorithm_name} Silhouette Score: {silhouette_avg:.4f}")
K-Means Silhouette Score: 0.4599
Hierarchical Agglomerative Silhouette Score: 0.4467
DBSCAN Silhouette Score: 0.5217
```

```
# Visualize the clustering results in a 2D plot
plt.figure(figsize=(8, 5))
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=labels, cmap='viridis', edgecolor='k', s=50)
plt.title(f'(algorithm_name) Clustering Results (Silhouette Score: {silhouette_avg:.4f})')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()
```

